

# “RATE MY BREW”

“INTRODUCTION TO DATA SCIENCE”

FINAL PROJECT

BY REMEZ SAVOY AND BAR NAHMUKA



## RESEARCH QUESTION

Can highly rated beers be predicted?




# STAGES

- Data Acquisition
- Data Cleaning
- EDA
- Machine Learning



# DATA ACQUISITION

- Scraping drizly.com
- Getting Beers ratings, countries and other beer properties






# SCRAPING DRIZLY

- Gathering all beer data from drizly.com using Selenium

Work plan: Entering the beer browsing section at Drizly, and going through all beers and scraping them to a CSV file.

# ENTERING THE BEER BROWSING SECTION

[Beer](#)[Wine](#)[Liquor](#)[Extras](#)[Corporate](#)[Gifts](#)


### Beer Styles


- IPA
- Hard Seltzer
- Lager
- Light Beer
- Ale
- Stout

### Trending

- Sour Beer
- Belgian-Style Ale
- Cider
- New England / Hazy IPA
- [Shop All Beer](#)

### Beer Brands

[Beer](#)[Wine](#)[Liquor](#)[Extras](#)[Corporate](#)[Gifts](#)[Sign in](#)[Create account](#)









[Shop now](#)

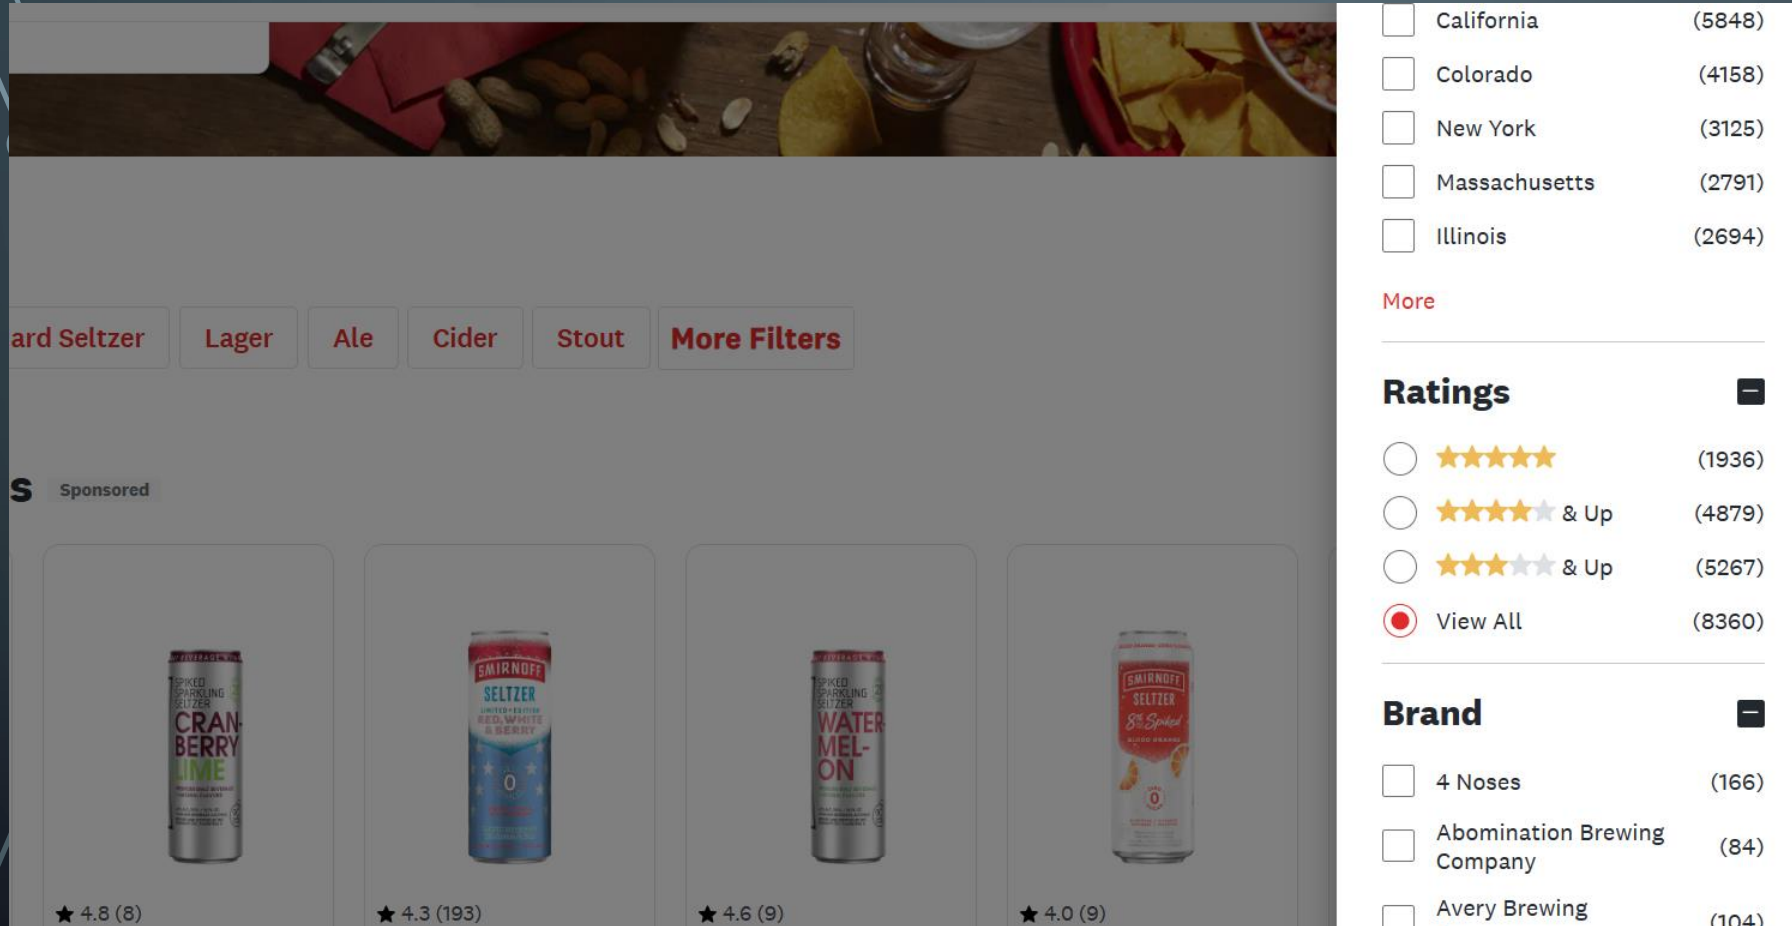
## Beer

[IPA](#)[Craft Beer](#)[Hard Seltzer](#)[Lager](#)[Ale](#)[Cider](#)[Stout](#)[More Filters](#)

### Beachin' beers Sponsored



# FILTERING RESULTS



The screenshot displays a beer filtering interface. At the top, a banner image shows a wooden surface with snacks like nuts and chips. Below the banner, a row of filter buttons includes 'ard Seltzer', 'Lager', 'Ale', 'Cider', 'Stout', and 'More Filters'. A 'Sponsored' label is visible on the left. The main product grid shows four beer cans: 'CRAN-BERRY LIME', 'SMIRNOFF SELTZER LIMITED EDITION RED, WHITE & BERRY', 'WATER-MEL-ON', and 'SMIRNOFF SELTZER 8th Spiked'. Each product has a star rating and a count in parentheses. On the right, a sidebar contains filters for 'Location' (California, Colorado, New York, Massachusetts, Illinois), 'Ratings' (5 stars, 4 stars & up, 3 stars & up, View All), and 'Brand' (4 Noses, Abomination Brewing Company, Avery Brewing).

Location	Count
<input type="checkbox"/> California	(5848)
<input type="checkbox"/> Colorado	(4158)
<input type="checkbox"/> New York	(3125)
<input type="checkbox"/> Massachusetts	(2791)
<input type="checkbox"/> Illinois	(2694)

[More](#)

Ratings	Count
<input type="radio"/> ★★★★★	(1936)
<input type="radio"/> ★★★★★ & Up	(4879)
<input type="radio"/> ★★★☆☆ & Up	(5267)
<input checked="" type="radio"/> View All	(8360)

Brand	Count
<input type="checkbox"/> 4 Noses	(166)
<input type="checkbox"/> Abomination Brewing Company	(84)
<input type="checkbox"/> Avery Brewing	(104)

Product	Rating	Count
CRAN-BERRY LIME	★ 4.8	(8)
SMIRNOFF SELTZER LIMITED EDITION RED, WHITE & BERRY	★ 4.3	(193)
WATER-MEL-ON	★ 4.6	(9)
SMIRNOFF SELTZER 8th Spiked	★ 4.0	(9)

We filtered the results to show beers rated 3+



# NOW WE GOT TO THIS SCREEN

The screenshot displays a grid of beer products on a website. The first row contains four items, each marked as an 'Ad'. The second row contains four items, with the first three marked as 'Best Seller'. The pagination bar at the bottom shows page 1 is active, with a total of 221 items across 22 pages.


Product Name	Rating (Stars)	Reviews (Count)	Price Range	Special Label
Budweiser	4.6	2,406	\$5.01 - \$59.99	Ad
Michelob Ultra	4.7	4,203	\$5.09 - \$59.99	Ad
Michelob Ultra Organic Seltzer - Classic Collection	4.3	62	\$8.79 - \$39.99	Ad
Coors Light American Lager Beer	4.7	4,894	\$5.03 - \$59.99	
Corona Extra	4.7	7,408	\$5.19 - \$59.99	Best Seller
Lite (A Time Pilsner Beer)	4.7	4,579	\$5.04 - \$59.99	Best Seller
Heineken	4.7	2,304		Best Seller

1 2 3 4 ... 221 >

And then we went through all the pages and scraped the beer data



# EXAMPLE FOR BEER PAGE

[Beer](#) [Wine](#) [Liquor](#) [Extras](#) [Corporate](#) [Gifts](#)

[Sign in](#) | [Create account](#) [Cart](#)

Enter A Delivery Address

[Drizly](#) / [Beer](#) / [Lager](#) / [Pale Lager](#) / [American-Style Lager](#)



## Budweiser

Best Seller

★★★★★

4.6 242 Reviews

American-Style Lager / 5% ABV / Missouri, United States

\*Packaging may vary

### You might also like

Enter a delivery address

Check availability

☒ 12 pack - from \$18.35

☐ 6 pack - from \$9.99

[View more sizes](#)

# EACH BEER CONTAINS THE FOLLOWING DETAILS

## Product details


Category	American-Style Lager
Region	Missouri, United States
ABV	5%
IBU	12
Tasting Notes	Balanced, Grainy, Round
Food Pairing	Beef, Chicken, Nuts
Suggested Glassware	Pint Glass
Suggested Serving Temperature	35-40° F

[Read Less](#)

These are the details we scraped for each beer



# SCRAPER CODE



```
def gather_data_from_item1(driver):
    #name = None
    try:
        name = driver.find_element(By.CLASS_NAME, 'jQELS7x16epG8DDVkkBa').text
    except:
        return None
    try:
        rating = driver.find_element(By.CLASS_NAME, 'btEBjbUX5Q2GMZaxsd_e').text
    except:
        rating = np.nan
    try:
        string = driver.find_element(By.CLASS_NAME, 'Ze_X7NgmFBRwchWfVpkC').text
        reviews = re.findall(r'\d+', string)[0]
    except:
        reviews = np.nan
    element_texts = []
    line_text = []
    try:
        driver.find_element(By.CLASS_NAME, "rgI6wOnWkzQAcuODfjPg").click()
    except:
        pass
    time.sleep(1)
    for t in driver.find_elements(By.CLASS_NAME, 'eGvxSSyFJKtDAmMxUYsQ'):
        line = t.text
        line_text.append(line)
    category = np.nan
    region = np.nan
    type = np.nan
    ABV = np.nan
    IBU = np.nan
    Tasting_Notes = np.nan
    Food_Pairing = np.nan
    Suggested_Glassware = np.nan
    Suggested_Serving_Temperature = np.nan
    Calories_Per_Serving = np.nan
    Carbs_Per_Serving = np.nan
    Features = np.nan
```



```
for element in driver.find_elements(By.CLASS_NAME, 'B7w58_0hngVFP2eha4P9'):
    ele = element.text
    element_texts.append(ele)
for a in element_texts:
    if 'Category' in a:
        category = line_text[element_texts.index(a)]
    elif 'Type' in a:
        type = line_text[element_texts.index(a)]
    elif 'Region' in a:
        region = line_text[element_texts.index(a)]
    elif 'ABV' in a:
        ABV = line_text[element_texts.index(a)]
    elif 'IBU' in a:
        IBU = line_text[element_texts.index(a)]
    elif 'Tasting Notes' in a:
        Tasting_Notes = line_text[element_texts.index(a)]
    elif 'Food Pairing' in a:
        Food_Pairing = line_text[element_texts.index(a)]
    elif 'Suggested Glassware' in a:
        Suggested_Glassware = line_text[element_texts.index(a)]
    elif 'Suggested Serving Temperature' in a:
        Suggested_Serving_Temperature = line_text[element_texts.index(a)]
    elif 'Calories Per Serving' in a:
        Calories_Per_Serving = line_text[element_texts.index(a)]
    elif 'Carbs Per Serving' in a:
        Carbs_Per_Serving = line_text[element_texts.index(a)]
    elif 'Features' in a:
        Features = line_text[element_texts.index(a)]
try:
    price_box = driver.find_element(By.XPATH, '/html/body/div[5]/main/div/div[2]')
    price_temp = price_box.find_element(By.XPATH, '/html/body/div[5]/main/div/div[2]/div/fieldset/div/label[1]/span[2]/s
    matches = re.findall(r'\d+\.\d+', price_temp)
    price = float(matches[0])
except:
    price = np.nan
```



```
return {  
  'Name': name,  
  'Price': price,  
  'Rating': rating,  
  'Reviews': reviews,  
  'Category' : category,  
  'Region' : region,  
  'Type': type,  
  'ABV' : ABV,  
  'IBU' : IBU,  
  'Tasting_Notes' : Tasting_Notes,  
  'Food_Pairing' : Food_Pairing,  
  'Suggested_Glassware' : Suggested_Glassware,  
  'Suggested_Serving_Temperature' : Suggested_Serving_Temperature,  
  'Calories Per Serving (12 Oz)': Calories_Per_Serving,  
  'Carbs Per Serving (12 Oz)': Carbs_Per_Serving,  
  'Features' : Features  
}
```

# POPUP PROTECTION & SITE LINKING FROM SELENIUM



```
driver = uc.Chrome()

#driver = webdriver.Chrome()
driver.maximize_window()
driver.get("https://drizly.com/beer/c2/page1?r=3")
time.sleep(5)
driver.find_element(By.CLASS_NAME, "uYf9y6pN6bXzrDlVhFGb").click()
#time.sleep(5)
mouse = Controller()
mouse.position = (235,325)
mouse.move(234,-234)
df_list = []
num = 0
for i in range(220):
    elem_links_list = []
    container = driver.find_element(By.XPATH, "/html/body/div[5]/main/div[3]/div[1]/section[1]/ul")
    time.sleep(3)
    css_element = container.find_elements(By.CSS_SELECTOR, "*")

    for temp_elm in css_element:
        href = temp_elm.get_attribute('href')
        if href:
            elem_links_list.append(href)
    #print(elem_links_list)

    for t in elem_links_list:
        driver.get(t)
        time.sleep(5)
        sm_data = gather_data_from_item1(driver)
        if sm_data is None:
            continue
        print(sm_data)
        df_list.append(pd.DataFrame(sm_data, index=[num]))
        num += 1
    df = pd.concat(df_list, ignore_index=True)
    df.to_csv('beer1.csv', index=False)
```

[illegible]



# COLUMNS

- Name
- Price
- Rating
- Reviews
- Category
- Region
- Type
- ABV
- IBU
- Tasting Notes
- Food Pairing
- Suggested Glassware
- Suggested Serving Temperature
- Calories Per Serving (12 Oz)
- Carbs Per Serving (12 Oz)
- Features



# DATA CLEANING

- Handling missing values
- Removing '%' marks from data
- Converting all sort of data from object form to integers
- Removing beers rated '0'
- Removing city names and keeping only countries for each beer
- Dealing with outliers
- Dropping duplicates by 'Name'



# HANDLING MISSING VALUES

```
df_copy['Suggested_Serving_Temperature (F)'].fillna(avg_sst, inplace=True)
df_copy['Type'].fillna('Craft', inplace=True)
df_copy['Rating'].fillna(0, inplace=True)
df_copy['Region'].fillna('Unknown', inplace=True)
df_copy['ABV'].fillna(avg_abv, inplace=True)
df_copy['IBU'].fillna(avg_ibu, inplace=True)
df_copy['Category'].fillna('Unknown', inplace=True)
df_copy['Reviews'].fillna(0, inplace=True)
df_copy['Features'].fillna('None', inplace=True)
df_copy['Suggested_Glassware'].fillna('Unknown', inplace=True)
df_copy['Calories Per Serving (12 Oz)'].fillna(avg_cps, inplace=True)
df_copy['Carbs Per Serving (12 Oz)'].fillna(avg_carps, inplace=True)
df_copy['Food_Pairing'].fillna('Unknown', inplace=True)
df_copy['Tasting_Notes'].fillna('No special notes', inplace=True)
df_copy['Price'].fillna(avg_p, inplace=True)
```

# REMOVING '%' MARKS FROM DATA

```
# Delete F
df_copy.loc[df_copy['Suggested_Serving_Temperature'].str.contains('-', na=False), 'Suggested_Serving_Temperature'] =
df_copy['Suggested_Serving_Temperature'].str.split('-').str[1].str.strip()
df_copy['Suggested_Serving_Temperature'] = df_copy['Suggested_Serving_Temperature'].str.replace('° F', '')
#convert to float
df_copy['Suggested_Serving_Temperature'] = df_copy['Suggested_Serving_Temperature'].astype(float)
#change column name
df_copy.rename(columns={'Suggested_Serving_Temperature': 'Suggested_Serving_Temperature (F)'}, inplace=True)

#ABV DELETE %
df_copy['ABV'] = df_copy['ABV'].str.replace('%', '')
df_copy['ABV'] = df_copy['ABV'].astype(float)
```

# REMOVING CITY NAMES AND KEEPING ONLY COUNTRIES FOR EACH BEER

```
# Delete cities  
df_copy.loc[df_copy['Region'].str.contains(',', na=False), 'Region'] = df_copy['Region'].str.split(',').str[1].str.strip()
```

# REMOVING BEERS RATED '0'

```
# Delete rating 0  
df_copy = df_copy[df_copy['Rating'] != 0.0]
```

# DROP DUPLICATES BY NAME

```
[4]: # Drop duplicates by Name  
df_copy.drop_duplicates(subset = ['Name'], inplace=True)
```

# AFTER DATA CLEANING

	Name	Price	Rating	Reviews	Category	Region	Type	ABV	IBU	Tasting_Notes	Food_Pairing	Suggested_Glassware	Suggested_Se
0	Budweiser	18.69	4.6	242.0	American-Style Lager	United States	Craft	5.0	12.0	Balanced, Grainy, Round	Beef, Chicken, Nuts	Pint Glass	
1	Heineken Silver Lager	19.99	4.6	5.0	Lager	Netherlands	Craft	4.0	5.0	No special notesbeer_updated	Unknown	Unknown	
2	Michelob Ultra Organic Seltzer - Classic Colle...	19.99	4.4	3.0	Hard Seltzer	United States	Variety Pack	4.0	35.0	No special notesbeer_updated	Unknown	Snifter/Goblet/Chalice	
3	Coors Light American Lager Beer	14.59	4.7	312.0	Light Lager	United States	Craft	4.2	10.0	Crisp, Dry, Light, Smooth	Nuts, Chicken	Pint Glass	
4	Corona Extra Lager Mexican Beer	19.99	4.7	2930.0	Pilsner	Mexico	Craft	4.6	18.0	Dry, Grainy, Light, Neutral	Nuts, Chicken	Pint Glass	
...	...	...	...	...	...	...	...	...	...	...	...	...	...
2404	Aslin Power Moves IPA	14.99	3.0	1.0	Imperial / Double IPA	United States	Craft	5.5	35.0	No special notesbeer_updated	Unknown	Snifter/Goblet/Chalice	
2405	Bronx Brewery Boogie Down Set	24.34	5.0	1.0	Variety Pack Beer	United States	Craft, Variety Pack	6.3	59.0	No special notesbeer_updated	Unknown	Pint Glass, Stein/Pub Mug, Snifter/Goblet/Chalice	
2406	Three Taverns Prince Of Pilsen	12.00	5.0	1.0	Pilsner	United States	Craft	5.0	35.0	No special notesbeer_updated	Unknown	Pilsner Glass	

2409 rows × 16 columns



```
df_copy.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 2409 entries, 0 to 4154
```

```
Data columns (total 16 columns):
```


#	Column	Non-Null Count	Dtype
0	Name	2409 non-null	object
1	Price	2409 non-null	float64
2	Rating	2409 non-null	float64
3	Reviews	2409 non-null	float64
4	Category	2409 non-null	object
5	Region	2409 non-null	object
6	Type	2409 non-null	object
7	ABV	2409 non-null	float64
8	IBU	2409 non-null	float64
9	Tasting_Notes	2409 non-null	object
10	Food_Pairing	2409 non-null	object
11	Suggested_Glassware	2409 non-null	object
12	Suggested_Serving_Temperature (F)	2409 non-null	float64
13	Calories Per Serving (12 Oz)	2409 non-null	float64
14	Carbs Per Serving (12 Oz)	2409 non-null	float64
15	Features	2409 non-null	object

```
dtypes: float64(8), object(8)
```

```
memory usage: 319.9+ KB
```



# EDA




# DEALING WITH OUTLIERS

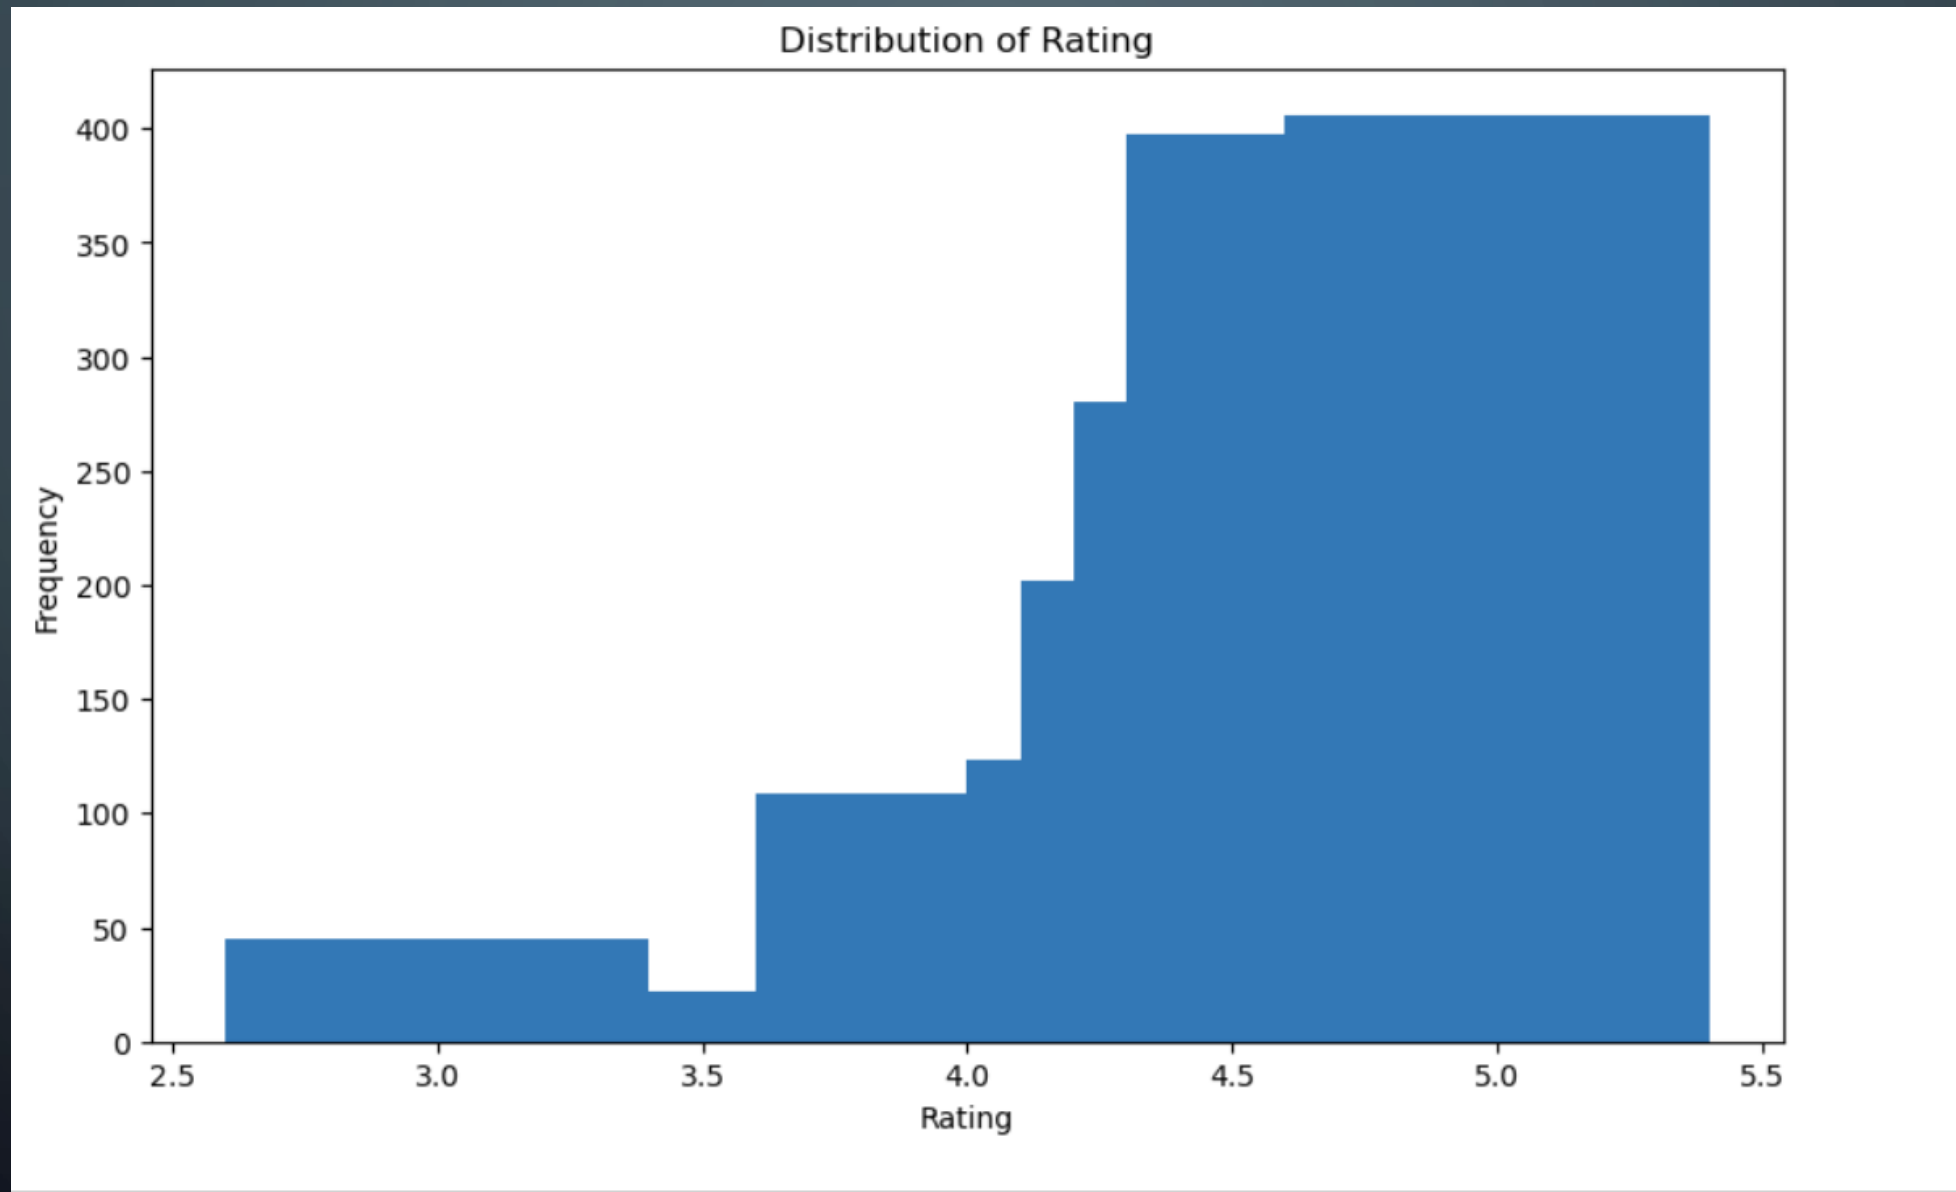
```
In [20]: def find_outliers_IQR(df):  
        # Convert 'Rating' column to numeric type  
        df['Rating'] = pd.to_numeric(df['Rating'], errors='coerce')  
  
        q1 = df['Rating'].quantile(0.25)  
        q3 = df['Rating'].quantile(0.75)  
        IQR = q3 - q1  
        outliers = df[((df['Rating'] < (q1 - 1.5 * IQR)) | (df['Rating'] > (q3 + 1.5 * IQR)))]  
        return outliers  
outliers = find_outliers_IQR(df_copy)  
  
print('Number of outliers: ' + str(len(outliers)))  
print('Max outlier value: ' + str(outliers['Rating'].max()))  
print('Min outlier value: ' + str(outliers['Rating'].min()))  
  
print(outliers)
```

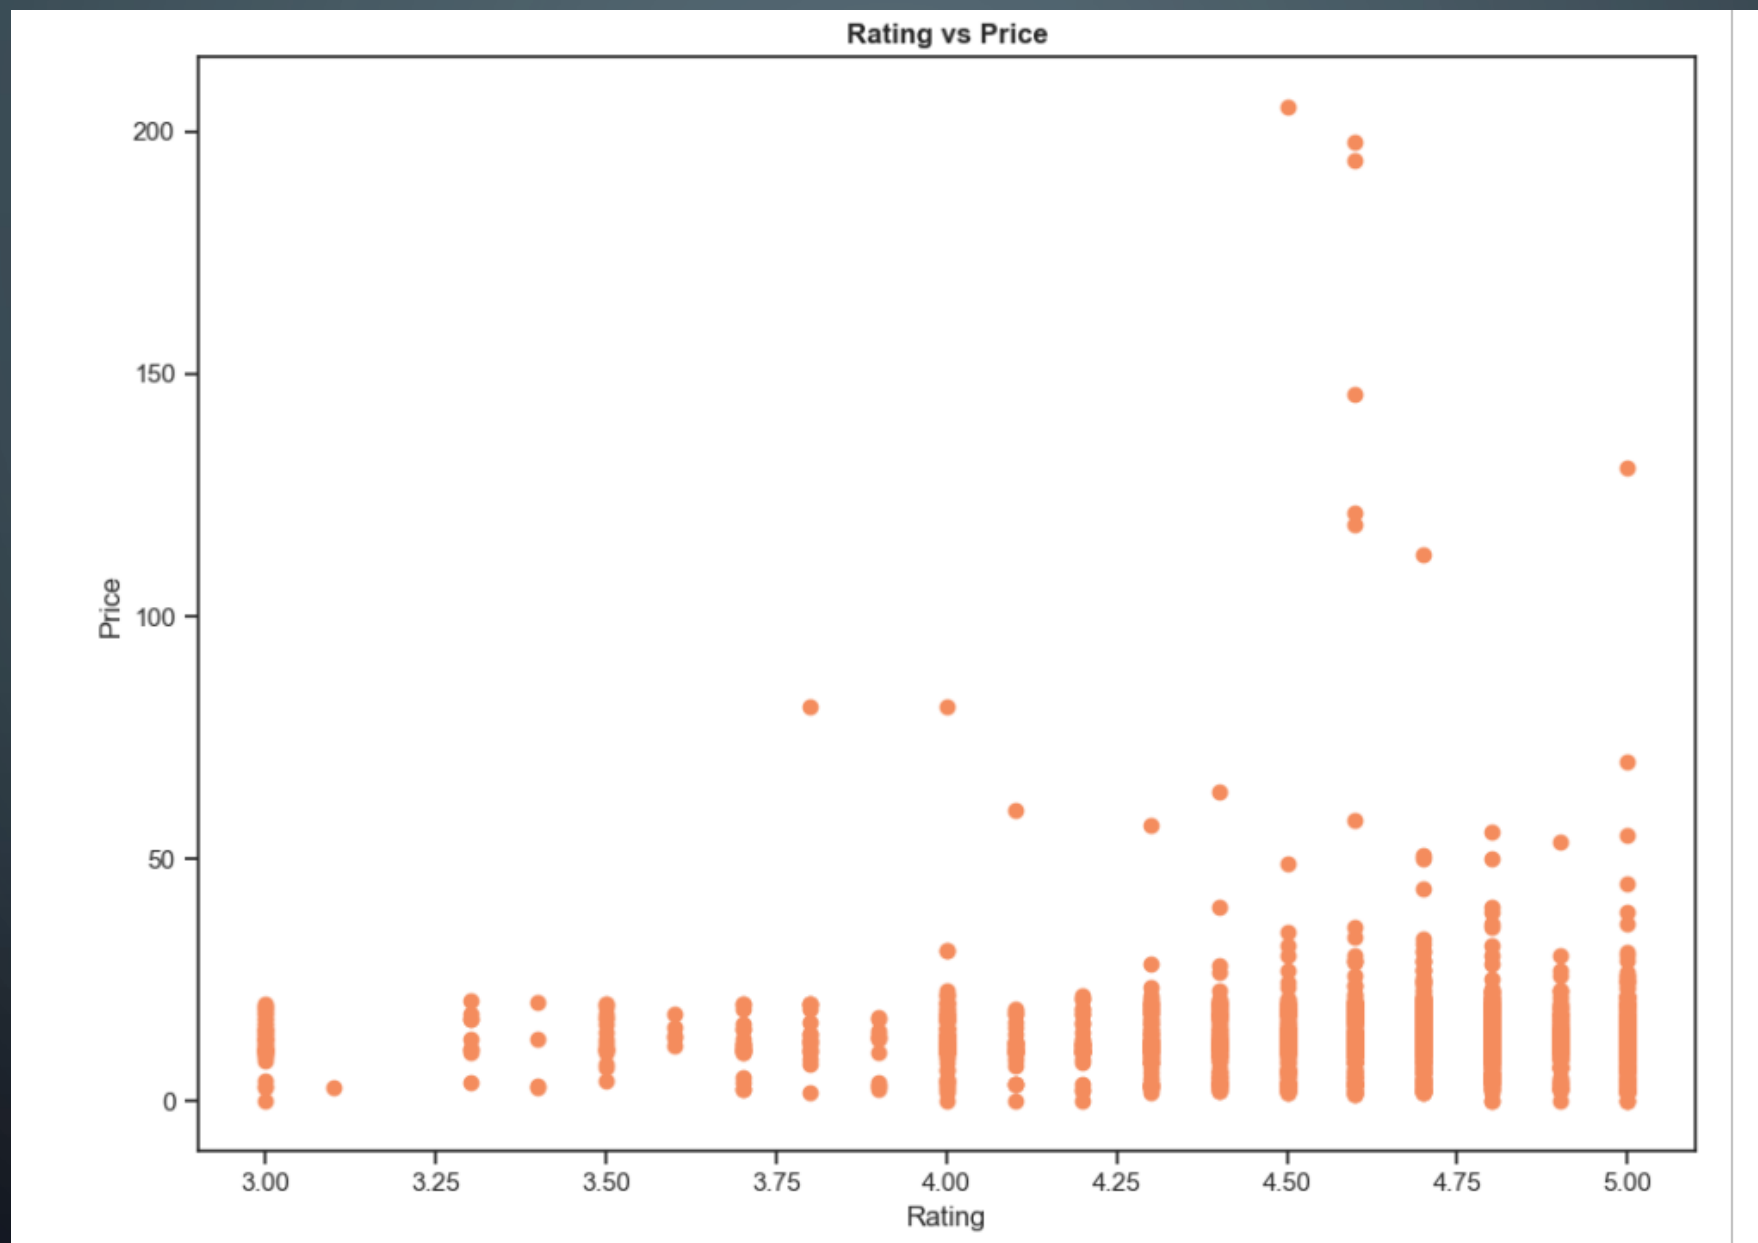
```
Number of outliers: 123  
Max outlier value: 3.8  
Min outlier value: 3.0
```


	Name	Price	Rating	\
63	Flying Embers Wild Berry Hard Kombucha	15.39	3.6	
254	Happy Dad Hard Seltzer Death Row Records Grape	19.99	3.7	
312	Modelo Chelada Sandia Picante Mexican Import F...	3.99	3.3	
693	Founders All Day Haze, Session Hazy IPA Beer	18.99	3.8	
709	Sierra Nevada Seasonal Oktoberfest	0.00	3.0	
...	...	...	...	
6021	Fireball X Lemonade	3.38	3.4	
6082	Pontoon Down With The Thickness	16.22	3.8	
6195	Right Proper Senate Beer Lager	13.99	3.8	
6228	Aslin Power Moves IPA	14.99	3.0	
6235	1911 Maple Bourbon Barrel Aged Hard Cider	14.99	3.0	

# DEALING WITH OUTLIERS











Cross-Tabulation: Rating vs Features




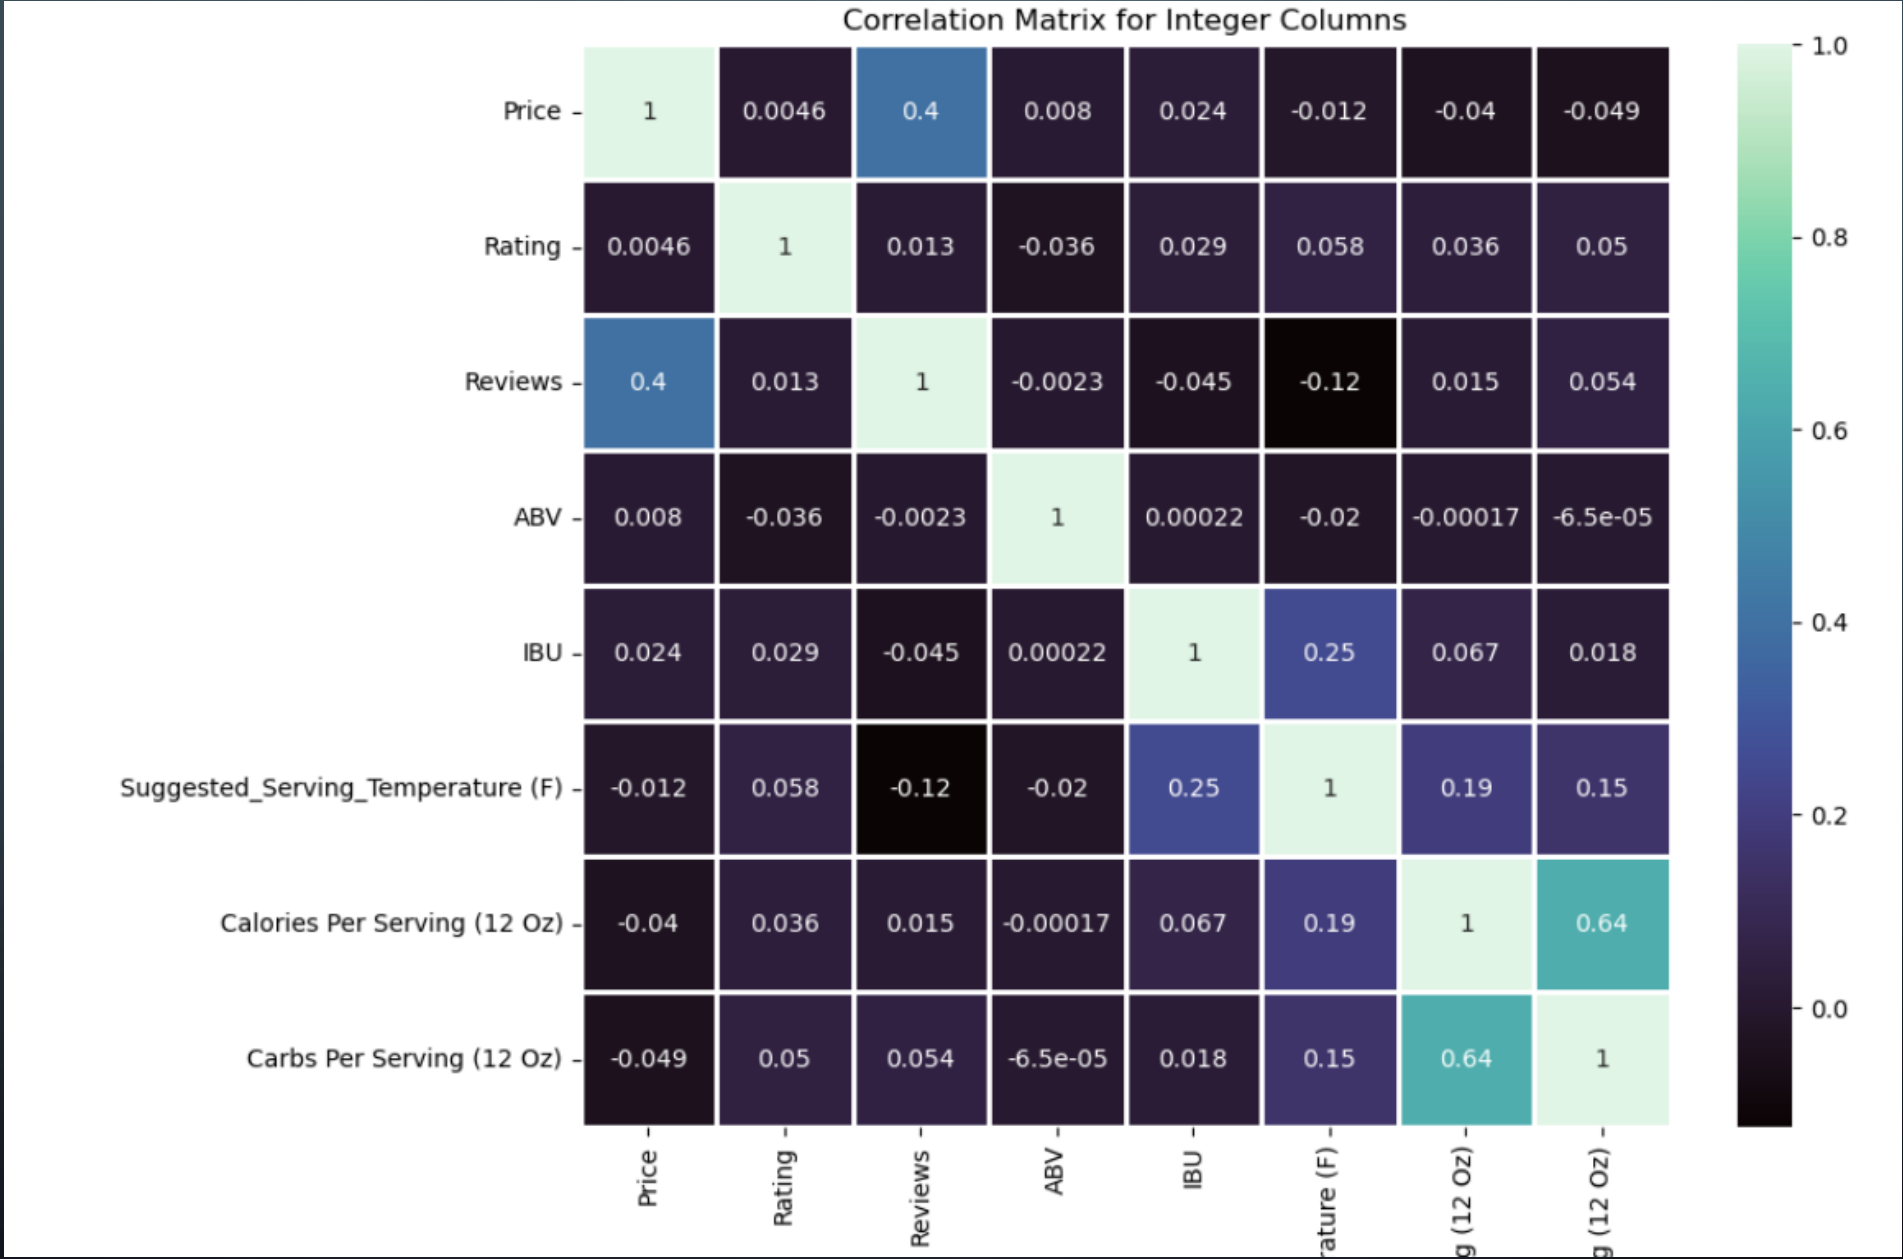



Comparison of Rating and Category

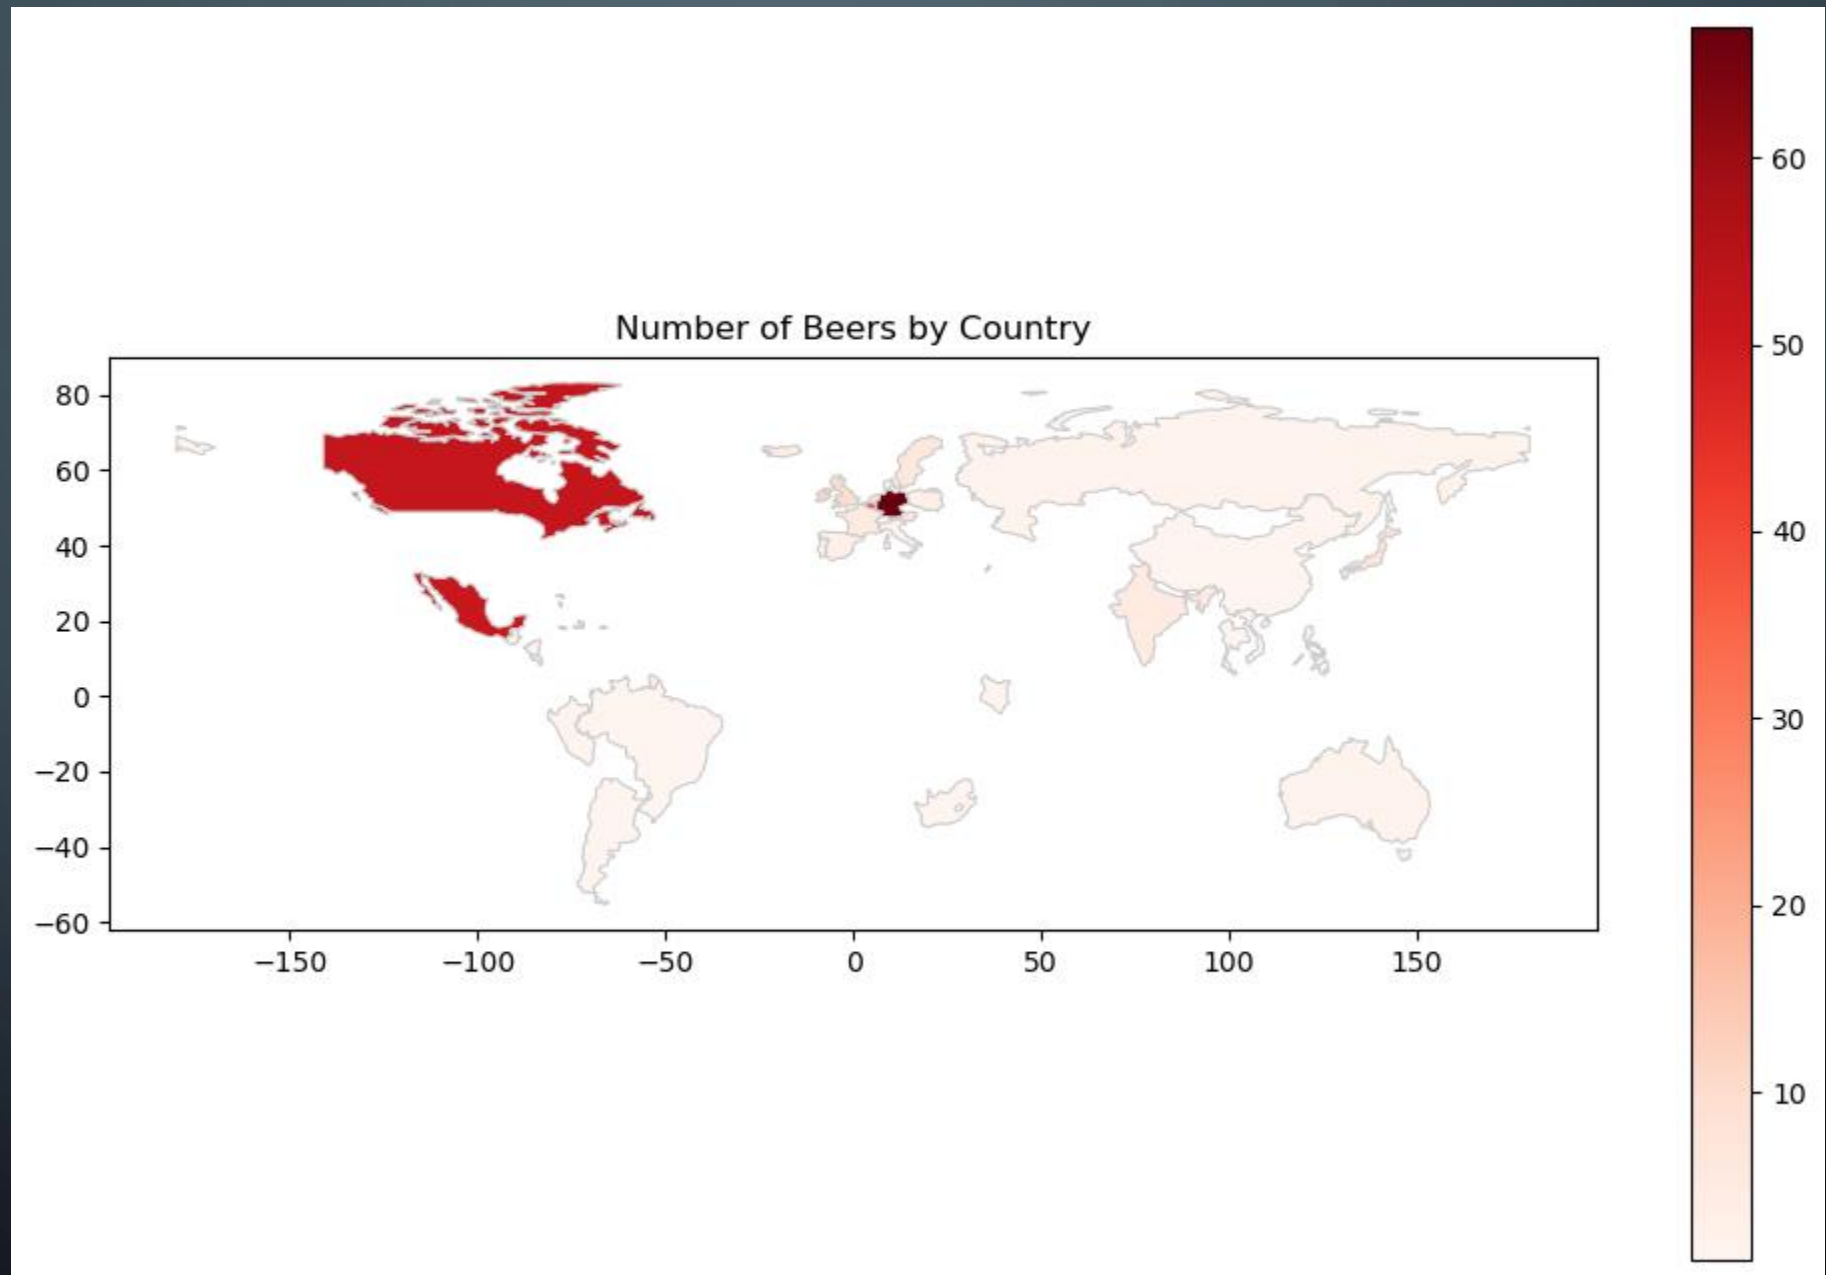


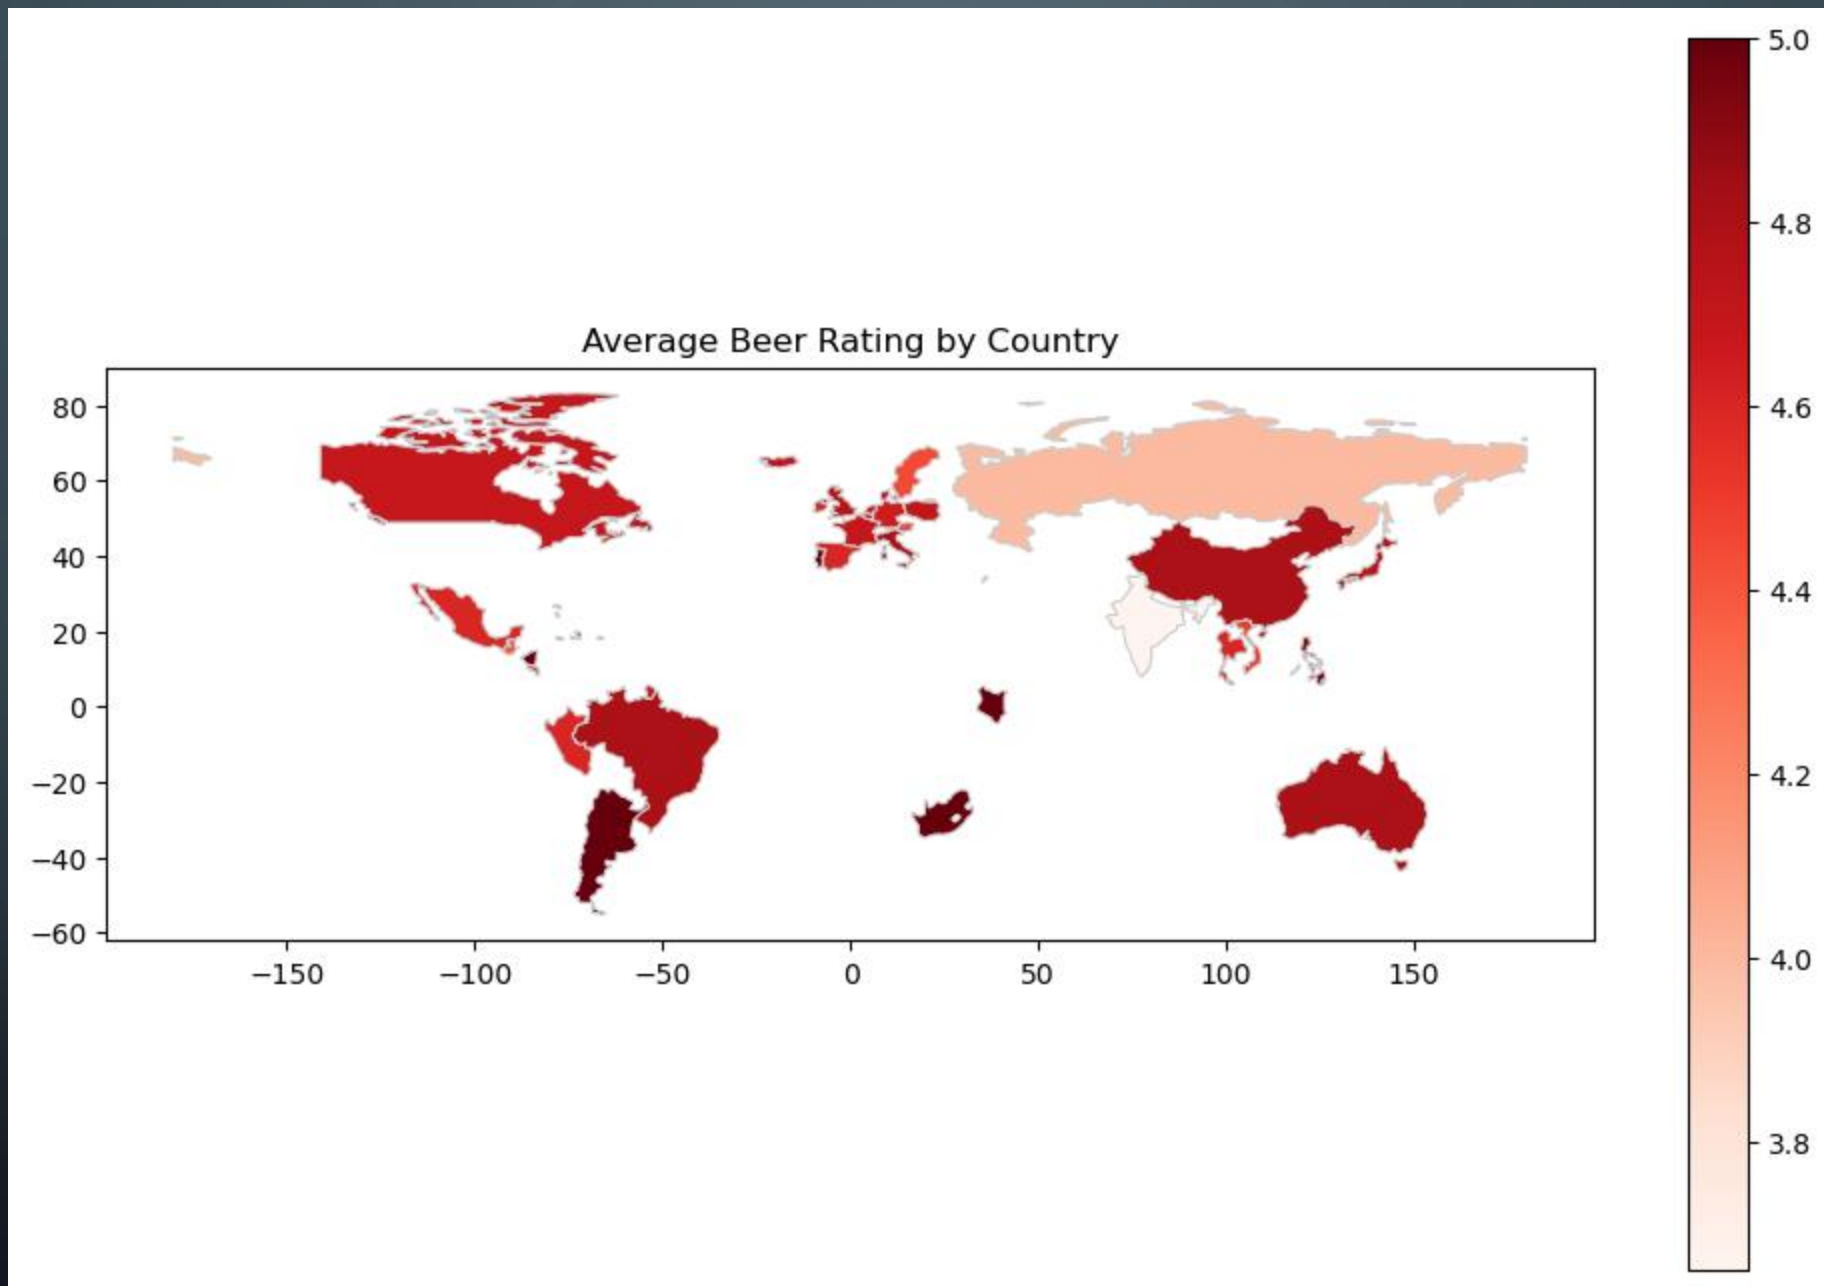
Comparison of Rating and Region

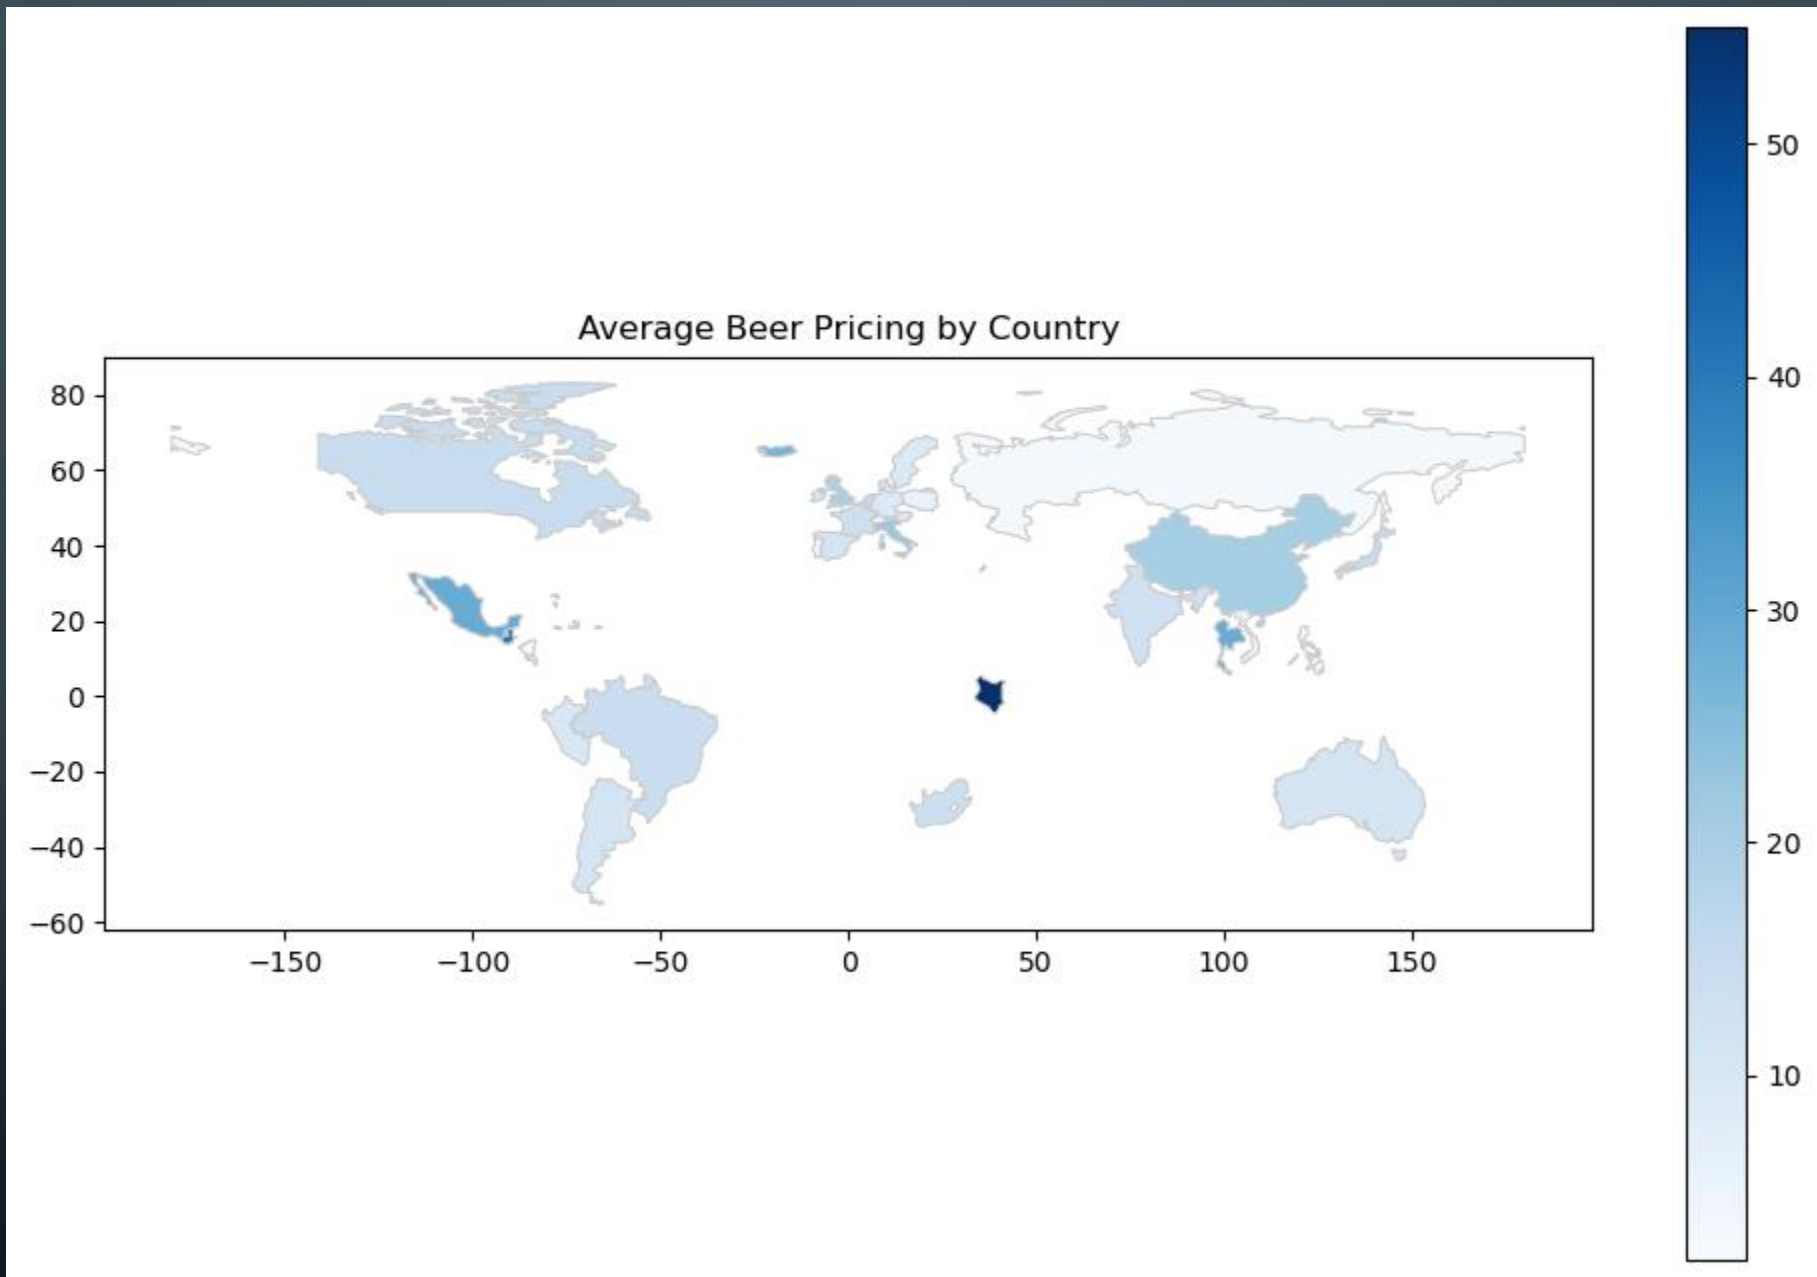














# MACHINE LEARNING

- Preparing the data
- Comparing different models
- Hyperparameter tuning
- Final evaluation



# PREPARING THE DATA

- Splitting data to train/test
- Scaling numeric features
- Min/max scaler features
- PCA



# PREPARING THE DATA – CODE / FIRST MODEL

```
def func(X,y,cols):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 40)
    #scaling the numeric columns
    scaler = MinMaxScaler()
    numeric_cols = X_train.select_dtypes(include = ['float64', 'int64']).columns
    X_train_numeric_scaled = scaler.fit_transform(X_train[numeric_cols])
    X_test_numeric_scaled = scaler.transform(X_test[numeric_cols])
    #transform it into dataframe
    X_train_numeric_scaled = pd.DataFrame(X_train_numeric_scaled, columns = numeric_cols, index = X_train.index)
    X_test_numeric_scaled = pd.DataFrame(X_test_numeric_scaled, columns = numeric_cols, index = X_test.index)

    df_new = df_copy.copy()
    #transform columns from categorical to binary categorical column
    for a in cols:
        X[a] = X[a].str.replace('/', ',')
        X_train_numeric_scaled = X_train_numeric_scaled.merge(X[a].str.get_dummies(sep = ',').loc[X_train.index,:], left_index=True, right_index=True)
        X_test_numeric_scaled = X_test_numeric_scaled.merge(X[a].str.get_dummies(sep = ',').loc[X_test.index,:], left_index=True, right_index=True)

    #Conducting PCA
    pca = PCA(0.9)
    X_train_PCA = pca.fit_transform(X_train_numeric_scaled)
    X_test_PCA = pca.transform(X_test_numeric_scaled)

    return X_train_PCA,y_train,X_test_PCA,y_test
```



# PREPARING THE DATA – CODE / FIRST MODEL

```
# Getting avg cross validation score for model
def cv_avg(model, X, y):
    kfold = KFold(n_splits=10, shuffle=True, random_state=42)
    scores = cross_val_score(model, X, y, cv=kfold, scoring='r2')
    return scores.mean()

# We will now try three methods that we learned about in the course
cols = ['Region', 'Type', 'Suggested_Glassware', 'Food_Pairing']
X_train_PCA, y_train, X_test_PCA, y_test = func(X, y, cols)
regression_model = LinearRegression()
linear_regression_avg = cv_avg(regression_model, X_train_PCA, y_train)
print("LinearRegression: " + str(linear_regression_avg))

svm_model = SVR()
svm_avg = cv_avg(svm_model, X_train_PCA, y_train)
print("SVR: " + str(svm_avg))

knn_model = KNeighborsRegressor()
knn_avg = cv_avg(knn_model, X_train_PCA, y_train)
print("KNeighborsRegressor: " + str(knn_avg))
```

## RESULTS

```
LinearRegression: 0.015336781723782267
SVR: -0.05296502473295228
KNeighborsRegressor: -0.1374952346744758
```

# PREPARING THE DATA – CODE / SECOND MODEL

```
# Second model
def create_regression_models(df_copy):
    df_copy = df_copy.copy()

    categorical_cols = ['Suggested_Glassware', 'Food_Pairing', 'Features']
    X = df_copy[['Price', 'Reviews', 'ABV', 'IBU']]
    y = df_copy['Rating']
    # Initialize the LabelEncoder
    encoder = LabelEncoder()

    for col in categorical_cols:
        df_copy[col] = encoder.fit_transform(df_copy[col])

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

    # Scaling numeric features
    scaler = MinMaxScaler()
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)

    # Initialize and train the models
    models = [
        LinearRegression(),
        svm.SVR(),
        KNeighborsRegressor()
    ]

    r_score = []
    pca = PCA(0.9)
    X_train_PCA = pca.fit_transform(X_train_scaled)
    X_test_PCA = pca.transform(X_test_scaled)
```



# PREPARING THE DATA – CODE / SECOND MODEL

```
r_score = []
pca = PCA(0.9)
X_train_PCA = pca.fit_transform(X_train_scaled)
X_test_PCA = pca.transform(X_test_scaled)

for model in models:
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)
    r2 = r2_score(y_test, y_pred)
    r_score.append(r2)
return r_score

r_score = create_regression_models(df_copy)

print("LinearRegression: " + str(r_score[0]))
print("SVR: " + str(r_score[1]))
print("KNeighborsRegressor: " + str(r_score[2]))
```

## RESULTS

```
LinearRegression: 0.001955039331048347
SVR: -0.046273346442465835
KNeighborsRegressor: -0.24771261136421296
```

# CONCLUSION

After examining the first model and the second model, we reached a better result in the first model than in the second model. But we did not reach a sufficient result in order to predict our research question and therefore our conclusion is that it is not possible way to predict the success of a beer using ratings of other beers.

The best prediction we've accomplished:

LinearRegression: 0.015336781723782267

